



PHASE 1 causal construction

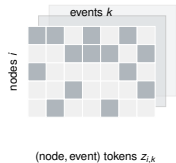
1. Edge stream



2. Evidence pool

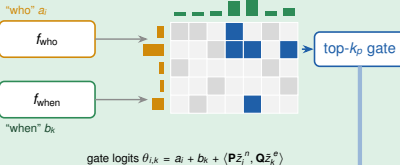


3. Tokens



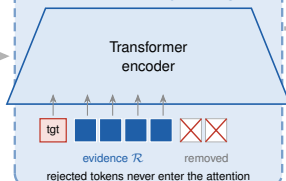
PHASE 2 selection

4. Evidence selector



PHASE 3 scoring

5. Evidence-only scoring



Anomaly score
 $s(e^*)$

Explanation
evidence $\mathcal{R}$
who α , when β

training only: the rejected tokens go through the same encoder;
their score s_c is pushed to the prior by $\mathcal{L}_{\text{nec}}$